

Component-based Software Development

Java Persistence APIs & Entity Manager

**Dr. Vinod Dubey
SWE 645
George Mason University**

Acknowledgement/Reference

<https://docs.oracle.com/javaee/6/tutorial/doc/bnbpz.html>

&

<http://www.vogella.de/articles/JavaPersistenceAPI/article.html>

Outline

- **Data Persistence**
- **Java Persistence API**
- **JPA Entities**
- **Persistence Unit**
- **Persistent Context**
- **Entity Manager**
- **Examples using JPA**

Data Persistence

- Mechanism to store application data into a persistence storage such as database

Java Object

Person

- id
- firstNme
- lastName

Table

PERSON_ID	FIRST_NAME	LAST_NAME
1	Greg	Martin
2	John	Doe
3	George	Smith

Data Persistence

- Mechanism to store application data into a persistence storage such as database
- Approach so far: **JDBC/SQL**

Java Object

Person

- id
- firstNme
- lastName

Table

PERSON_ID	FIRST_NAME	LAST_NAME
1	Greg	Martin
2	John	Doe
3	George	Smith

Data Persistence

- Mechanism to store application data into a persistence storage such as database
- Approach so far: **JDBC/SQL**
- New Approach: **O/R Mapping**
 - Object/Relational Mapping

Java Object

Person


- id
- firstName
- lastName

Table

PERSON_ID	FIRST_NAME	LAST_NAME
1	Greg	Martin
2	John	Doe
3	George	Smith

Object Relation Mapping (ORM)

The Java Persistence API utilizes the O/R mapping design.



Here a single JavaBean is mapped to a database record. This type of design allows for data manipulation using Objects.

When to Use ORM

- There is a **natural mapping** between **objects** and database **tables**
- Most create/update/delete operations work with individual objects
- You don't want to worry about the mechanics of the database, and vendor specific SQL

When NOT to Use ORM

- There is **no easy way to map objects to database tables**
- When **not working with a Relational Database**
- Data access is predominant through large data sets, batch operation

Goal

The Goal is to learn:

- How to **use** Java Persistence APIs (**JPA**) **annotations** to **configure** Java classes (POJOs) into **JPA entities** to provide object to relational mapping
- How to use **Persistence APIs** (EntityManager) to **implement CRUD** operations

JPA Programming Model

- **JPA Entities** are plain **Java classes** (POJOs) that are **mapped to** relational database table
 - Uses Java Persistence API (JPA) annotations
- **Entity Manager** – is used **to create, query, remove, and update** JPA entities in your java application
 - Which in turn updates records into the database table

JPA Programming Model

- **JPA is a Java specification that provides Object to Relational mapping**
 - It define **syntax** to configure a **java class** into a **JPA entity**, capture **primary** and **foreign key constraints** etc.
 - Provides **capability** to **create, read, update, and delete** database **rows** using JPA entities
 - Provides support for **transactions** and other similar capabilities required by an application accessing a database
 - Uses persistent metadata via **annotations** in the Java class

Java Persistence API (JPA)

- **Helps manage persistence of application data in Java Enterprise Edition (EE) and Java Standard Edition (SE) applications.**
 - Allows you to **directly work with Java classes** instead of database modules or SQL
- **Reference implementation of JPA: EclipseLink**
 - Other popular implementations: **Hibernate**, **Apache OpenJPA**

JPA Entity Classes

- **Two Bare Minimum Requirements to specify a JPA entity**
 - The **class** must be **annotated with** the **@Entity**
 - `javax.persistence.Entity`
 - Must define a **primary key** using **@Id** annotation
- **JPA entity class (POJO) with a no-arg public constructor is used to define the mapping with one or more relational database tables**
 - The instance variables follow **JavaBeans-style properties** and represent the **persistent state** of the entity

JPA Entity Classes: Standard POJOs with certain requirements

● Other Requirements

- The class must have a **public** or **protected**, no-argument constructor
 - The class may have other constructors.
- The **class must not** be declared **final**
 - No methods or persistent instance variables must be declared final
- If an entity instance is **passed by value** as a detached object, such as through a session bean's remote business interface, the class must implement **the Serializable interface**
- Entities may **extend** both **entity** and **non-entity classes**, and non-entity classes may extend entity classes
- **Persistent instance variables** must be **declared private**, protected, or package-private
 - Can be accessed directly only by the entity class's methods
 - **Clients** must **access** the **entity's state** through **accessor** or **business methods**

JPA Entity Classes (Contd.)

- By default, **table name** in the database **correspond to the Entity class name**
 - Default table name **can be changed** using `@Table(name="NEW_TABLE_NAME")`
- By default, **entity fields names are mapped to columns with same name**
 - Default names **can be changed** using `@Column` annotation, e.g., `@Column (name="newColumnName")`;

JPA Entity Classes (Contd.)

- **An Entity class must define a primary key**
 - @Id annotation
 - Can be auto generated using @GeneratedValue annotation
- **JPA persists all fields by default**
 - Fields that should not be persisted should be annotated with @Transient
- **Clients of entity beans make use of EntityManager to persist and manipulate data into database**
- **An instance of an Entity class represent a row in the table**

Persistent Field Types

- The **persistent fields** or properties of an entity may be of the following **types**:
 - Java primitive types,
 - `java.lang.String`,
 - `java.math.BigInteger`,
 - `java.math.BigDecimal`,
 - `java.util.Date`,
 - `java.util.Calendar`,
 - `java.sql.Date`,
 - `java.sql.Time`,
 - `java.sql.Timestamp`,
 - `byte[]`, `Byte[]`, `char[]`,
 - `Character[]`, enums and
 - other Java serializable types,
 - entity types, collections of entity types,
 - embeddable classes, and collections of basic and embeddable classes.
 - The **@Temporal** annotation may be specified on fields of type `java.util.Date` and `java.util.Calendar` to specify the temporal type of the field, e.g.,
`@Temporal(TemporalType.DATE)`
`private Date startDate; //puts date for startDate and not default timestamp`

Persistent Field Types

- For **String data**, hibernate, by default, **creates varchar(255)**
- If a field need to store **large amount of text**, you can **use @Lob** (large object) annotation,
 - which creates CLOB or BLOB data types and thus not restricted to 255 characters
- For example,
 @Lob
 private String comments;

JPA Entity – An example

- Two bare minimum annotations to configure a POJO as a JPA entity: **@Entity** and **@Id**

```
import javax.persistence.*;

@Entity
public class Customer implements java.io.Serializable{
    @Id
    private long id;
    private String firstName;
    private String lastName;

    public long getId() { return id; }
    public void setId(long id) { this.id = id; }

    public String getFirstName() {return firstName;}
    public void setFirstName(String fn) {this.firstName = fn;}

    public String getLastName() {return lastName;}
    public void setLastName(String ln) {this.lastName = ln;}
}
```

JPA Entity – An example

```
@Entity
@Table(name = "CUST_TBL")
public class Customer implements java.io.Serializable{
    private long id;
    private String firstName;
    private String lastName;

    @Id
    @Column(name="CUST_ID", nullable=false, columnDefinition="long")
    public long getId() { return id; }
    public void setId(long id) { this.id = id; }

    @Column(name="FIRST_NAME", length=255, nullable=false)
    public String getFirstName() {return firstName;}
    public void setFirstName(String fn) {this.firstName = fn;}

    @Column(name="LAST_NAME", length=255, nullable=false)
    public String getLastName() {return lastName;}
    public void setLastName(String ln) {this.lastName = ln;}

}
```

JPA Entity – An example

```
@Entity
@Table(name = "CUST_TBL")
public class Customer implements java.io.Serializable{
    private long id;
    private String firstName;
    private String lastName;

    @Id
    @Column(name="CUST_ID", nullable=false, columnDefinition="long")
    public long getId() { return id; }
    public void setId(long id) { this.id = id; }

    @Column(name="FIRST_NAME", length=255, nullable=false)
    public String getFirstName() {return firstName;}
    public void setFirstName(String fn) {this.firstName = fn;}

    @Column(name="LAST_NAME", length=255, nullable=false)
    public String getLastName() {return lastName;}
    public void setLastName(String ln) {this.lastName = ln;}
```

When @column annotations are on top of getters instead of fieldnames, the values inserted into tables are picked up from what is returned from getters

Primary Key

- A **primary key** is the **identity** of a given entity bean
 - Every entity bean must have a primary key, and it must be unique
- **Primary keys can map to one (or more) properties** and must map to one of the following **types**:
 - Any Java **primitive type**
 - `java.lang.String`, or
 - a **primary-key class** composed of primitives and/ or strings

Primary Key

- **Natural Key**

- Sometimes the primary **key** is **made up of real data that has real business use**, normally referred to as **natural keys**
- For example, if every customer has to have a **unique email id**, which can be **used as primary key**

- **Surrogate Key**

- **No business use**, but **acts as a primary key**
- For example, if no column is unique, you can have another **column** that has **unique number** and saved as primary key
- The surrogate **key** is **generated at runtime** when a new record is inserted into a table.
- A **surrogate key** is typically a **numeric value**

Primary Key: **@GeneratedValue** annotation

- You can generate the primary key for your entity beans manually or have the persistence provider do it for you.
- Use **@GeneratedValue** annotation to tell hibernate to **automatically generate** a **value** for **customerId**
 - Each time you persist a Customer object – can use next value sequence
 - **@javax.persistence.GeneratedValue** annotation
- **Don't have** to have **setter** method for **customerId**
- **Example**

```
@id @GeneratedValue  
private int customerId;
```

Primary Key: @GeneratedValue annotation

- Can use **strategy** attribute of **@GeneratedValue** to specify primary key generation type
- Four valid options are:
 - public enum GenerationType { TABLE, SEQUENCE, IDENTITY, AUTO};
- The **GeneratorType.AUTO** strategy is the **most commonly** used configuration, and it is the **default**
- **Example**

```
@id @GeneratedValue  
private int customerId;
```

Multiple Mapping with `@SecondaryTable`

One logical entity that is stored in two different tables

```
@Entity
@Table(name = "CUSTOMER_TABLE")
@SecondaryTable(name="ADDRESS_TABLE",
    pkJoinColumns={
        @PrimaryKeyJoinColumn(name="ADDRESS_ID")
    })
```

An Embedded Value object

- A **POJO** (e.g., Address) class that **does not** have a **persistent identity of its own** and **exclusively belongs to another entity** (e.g., Student class.)
 - This class is **identified by @Embedded** on the field in the **entity class** and annotated with **@Embeddable** in the class definition:

```
@Entity
public class Student implements Serializable {
    @Id
    private int id;
    private String name;
    private String grade;
    @Embedded
    private Address address;

    @ElementCollection
    @CollectionTable("StudentCourse")
    List<Course> courses;

    //...
}
```

Tells Address is embedded in Student and will be part of Student table, i.e., Address fields will be columns in Student table.

Tells this object is embedded someplace else, don't create separate table

```
@Embeddable
public class Address {
    private String street;
    private String city;
    private String zip;
    //...
}
```

Saving Collection

- The **@ElementCollection** annotation signifies that a student's **courses** are listed in a **different table**
 - By default, the **table name** is derived by combining the name of the **owning class**, the string “_,” and the **field name** (e.g., **Student_courses**.
 - **@CollectionTable** can be used to **override** the default name of the table
 - **@AttributeOverrides** can be used to override the default column names.

```
@Entity
public class Student implements Serializable {
    @Id
    private int id;
    private String name;
    private String grade;
    @Embedded
    private Address address;

    @ElementCollection
    @CollectionTable("StudentCourse")
    List<Course> courses;

    //...
}
```

Especially useful when you do not know how many Course objects you want to use in Student entity.

@ElementCollection can also be used for embedded class

EntityManager

- **The entities are managed by the EntityManager interface**
 - An object of this type manages data flow between the program and the database
- **The `javax.persistence.EntityManager` is the interface for all DB actions:**
 - **Persist** POJO into database
 - **Create** queries
 - **Find** objects
 - Synchronize objects
 - Cache
 - Transaction support

EntityManager Methods

Method Name	Description
<code>persist</code>	Makes an entity instance managed and persistent
<code>find</code>	Find an entity instance by executing a query by primary key on the database
<code>contains</code>	Returns true if the entity instance is in the <code>PersistentContext</code>
<code>merge</code>	Merges the state of a given entity into the current persistent context
<code>remove</code>	Removes the entity instance
<code>flush</code>	Forces the synchronization of the database with entities in the persistent context
<code>refresh</code>	Refreshes the entity instances in the persistent context from the database

Persistence Unit

- A **persistence unit** defines a set of all entity classes that are **managed** by **EntityManager** instance in an application
 - This set of entity classes represents the data contained within a **single data store**
- Persistence units are **defined** by the **persistence.xml** configuration file
 - An example **persistence.xml** file:

```
<persistence>
  <persistence-unit name="OrderManagement">
    <description>This unit manages orders and customers. It does not rely on any vendor-specific features
    and can therefore be deployed to any persistence provider.
    </description>
    <jta-data-source>jdbc/MyOrderDB</jta-data-source>
    <jar-file>MyOrderApp.jar</jar-file>
    <class>com.widgets.Order</class>
    <class>com.widgets.Customer</class>
  </persistence-unit>
</persistence>
```


Other JPA interfaces/Classes

- **EntityManagerFactory** interface
 - An object of this type **creates** persistence unit's **EntityManager**
- **Persistence** class
 - A static method of this class **creates** the specified persistence unit's **EntityManagerFactory**

Persistent Context

- A **working copy** of a **persistent unit**
- A set of managed entity instances that exist in a particular data store
- **Defines** the **scope** under which particular entity instances are created, persisted, and removed
- Each **EntityManager** instance is **associated** with **a persistence context**

Container-Managed Entity Managers

- **The Java EE container manages the lifecycle of container-managed entity managers**
- **To obtain an EntityManager instance, inject the entity manager into the application component:**
@PersistenceContext EntityManager em;

Application-Managed Entity Managers

- With an application-managed entity manager, the **lifecycle** of **EntityManager** instances is **managed** by the **application**
 - Applications create **EntityManager** instances by using the **createEntityManager** method of **javax.persistence.EntityManagerFactory**

```
@PersistenceUnit EntityManagerFactory emf;  
EntityManager em = emf.createEntityManager();
```
 - You can also use the **following** to get an instance of **EntityManagerFactory**

```
javax.persistence.Persistence.createEntityManagerFactory(String persistenceUnitName)
```

Application-Managed Entity Managers

- Such applications need to **manually** gain **access** to the **JTA transaction manager** and add transaction demarcation information when performing entity operations
- The `javax.transaction.UserTransaction` interface defines methods to **begin**, **commit**, and **roll back** transactions

Application-Managed Entity Managers

- **Inject an instance of `UserTransaction` by creating an instance variable annotated with `@Resource`:**
 - `@Resource UserTransaction utx;`
 - To begin a transaction, call the `utx.begin()` method.
 - When all the entity operations are complete, call the `utx.commit()` method to commit the transaction.
 - The `utx.rollback()` method is used to roll back the current transaction.
- **Alternatively, you can use `EntityManager`'s `getTransaction()` method to call `begin()`, `commit()`, `rollback()` -- See an example later**

Application-Managed Entity Managers

- **The following example shows how to manage transactions in an application that uses an application-managed entity manager:**

```
@PersistenceContext
EntityManagerFactory emf;
EntityManager em;
@Resource
UserTransaction utx;
...
em = emf.createEntityManager();
try {
    utx.begin();
    em.persist(SomeEntity);
    em.merge(AnotherEntity);
    em.remove(ThirdEntity);
    utx.commit();
} catch (Exception e) {
    utx.rollback();
}
```

EntityManager Services

Persisting Entities

```
Customer cust = new Customer();  
cust.setFirstName("Greg");  
entityManager.persist(cust);
```

Finding Entities

```
find(Class clazz, Object primaryKey)  
getReference(Class clazz, Object primaryKey)  
createQuery(String qry)
```

Update Entities

```
Customer cust = entityManager.find(...);  
cust.setXXXX(...);
```

Merging Entities

Takes a detached object and merges changes back to persistent storage.

```
entityManager.merge(obj);
```


EntityManager Services

Remove Entities

```
Customer cust = entityManager.find(...);  
entityManager.remove(cust);
```

Refreshing Entities

Refreshes the state of an object from the persistent storage.

```
entityManager.refresh(obj)
```

contains() and clear()

contains(Object entity) – determine **whether** an object is managed by the entity manager.

clear() detach an object from the entity manager.

Flushing Data

flush() - **force synchronization** with persistent storage

FlushModeType - AUTO, COMMIT

JPA Implementation Example using EclipseLink

```
package de.vogella.jpa.simple.model;

import javax.persistence.Entity;
import javax.persistence.GeneratedValue;
import javax.persistence.GenerationType;
import javax.persistence.Id;

@Entity
public class Todo {
    @Id
    @GeneratedValue(strategy = GenerationType.IDENTITY)
    private Long id;
    private String summary;
    private String description;

    public String getSummary() {
        return summary;
    }

    public void setSummary(String summary) {
        this.summary = summary;
    }

    public String getDescription() {
        return description;
    }

    public void setDescription(String description) {
        this.description = description;
    }

    @Override
    public String toString() {
        return "Todo [summary=" + summary + ", description=" + description
            + "]";
    }
}
```

JPA Entity: Todo

}

4.2. Persistence Unit

Create a directory "META-INF" in your "src" folder and create the file "persistence.xml". This examples uses EclipseLink specific flags for example via the parameter "eclipseLink.ddl-generation" you specify that the database scheme will be automatically dropped and created.

```
<?xml version="1.0" encoding="UTF-8" ?>
<persistence xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
  xsi:schemaLocation="http://java.sun.com/xml/ns/persistence
http://java.sun.com/xml/ns/persistence/persistence_2_0.xsd"
  version="2.0" xmlns="http://java.sun.com/xml/ns/persistence">
  <persistence-unit name="todos" transaction-type="RESOURCE_LOCAL">
    <class>de.vogella.jpa.simple.model.Todo</class>
    <properties>
      <property name="javax.persistence.jdbc.driver"
value="org.apache.derby.jdbc.EmbeddedDriver" />
      <property name="javax.persistence.jdbc.url"
value="jdbc:derby:/home/vogella/databases/simpleDb;create=true" />
      <property name="javax.persistence.jdbc.user" value="test" />
      <property name="javax.persistence.jdbc.password" value="test" />

      <!-- EclipseLink should create the database schema automatically -->
      <property name="eclipseLink.ddl-generation" value="create-
tables" />
      <property name="eclipseLink.ddl-generation.output-mode"
value="database" />
    </properties>
  </persistence-unit>
</persistence>
```

```
package de.vogella.jpa.simple.main;

import java.util.List;

import javax.persistence.EntityManager;
import javax.persistence.EntityManagerFactory;
import javax.persistence.Persistence;
import javax.persistence.Query;

import de.vogella.jpa.simple.model.Todo;

public class Main {
    private static final String PERSISTENCE_UNIT_NAME = "todos";
    private static EntityManagerFactory factory;

    public static void main(String[] args) {
        factory = Persistence.createEntityManagerFactory(PERSISTENCE_UNIT_NAME);
        EntityManager em = factory.createEntityManager();
        // Read the existing entries and write to console
        Query q = em.createQuery("select t from Todo t");
        List<Todo> todoList = q.getResultList();
        for (Todo todo : todoList) {
            System.out.println(todo);
        }
        System.out.println("Size: " + todoList.size());

        // Create new todo
        em.getTransaction().begin();
        Todo todo = new Todo();
        todo.setSummary("This is a test");
        todo.setDescription("This is a test");
        em.persist(todo);
        em.getTransaction().commit();

        em.close();
    }
}
```

A client
managing
using En

JPA Providers

- **Hibernate**
- **EclipseLink -- reference implementation**
- **Apache OpenJPA**

JPA Providers

- **Hibernate**

- To use hibernate, please Download hibernate (zip file) www.hibernate.org/downloads and jar files to project build path/classpath and add them to lib folder of your application
- On download page, click on release bundle, e.g., 3.6.4.....zip and, which has
 - has lib folder which has jars needed by hibernate
 - You will need all jars in [lib/jpa](#) and [lib/required](#) folders
- Also, download JDBC driver (of database choice) and add it to your project's java build path
 - This configures hibernate to use JDBC driver (jar) to connect to database

JPA Providers

hibernate.org/orm/downloads/



ORM Search Validator OGM Tools Others

Blog Community Follow Us



Hibernate ORM Downloads

About

Downloads

Docs (5.1)

Docs (5.0)

Docs (4.3)

Docs (4.2)

Tooling

Envers

Paid support

Get Certified

FAQ

Roadmap

Contribute

Wiki

Issues

Security issue

Forum

Source code

CI

stable 5.1.0.Final

Interested in commercial support? Check out Red Hat's offering.

Releases

5.1.0.Final



2016-02-10

stable

Maven gav: org.hibernate:hibernate-core:5.1.0.Final

Entity joins, load-by-multiple-ids, association traversal in AuditQuery

[More on this release](#)

5.0.9.Final



2016-03-14

stable

Maven gav: org.hibernate:hibernate-core:5.0.9.Final

Improved bootstrapping, hibernate-java8, hibernate-spatial, Karaf support

[More on this release](#)

4.3.11.Final



2015-08-05

stable

Maven gav: org.hibernate:hibernate-core:4.3.11.Final

JPA 2.1 support

[More on this release](#)

Adding Hibernate/JPA jars to your classpath in eclipse

- Unzip the downloaded file. Your jars are inside lib folder, specifically inside **lib/required** and **lib/jpa**

Documents library
hibernate-release-5.1.0.Final

Name	Date modified	Type	Size
documentation			
lib			
changelog.txt			
hibernate_logo.gif			
lgpl.txt			

Documents library
lib

Name
optional
envers
java8
jpa
jpa-metamodel-generator
osgi
required

Documents library
required


Name	Date modified
hibernate-core-5.1.0.Final.jar	2/10/2016 10:21 AM
hibernate-commons-annotations-5.0.1.Final.jar	11/30/2015 9:22 AM
jandex-2.0.0.Final.jar	11/30/2015 9:22 AM
javassist-3.20.0-GA.jar	10/6/2015 7:29 PM
classmate-1.3.0.jar	10/6/2015 10:12 AM
jboss-logging-3.3.0.Final.jar	5/28/2015 12:35 PM
geronimo-jta_1.1_spec-1.1.1.jar	5/5/2015 11:26 AM
hibernate-jpa-2.1-api-1.0.0.Final.jar	4/28/2014 8:30 PM
antlr-2.7.7.jar	4/28/2014 8:30 PM
dom4j-1.6.1.jar	4/28/2014 8:28 PM

Documents library
jpa

Name	Date modified
hibernate-entitymanager-5.1.0.Final.jar	2/10/2016 10:21 AM

Download MySQL JDBC driver and add it to Java Build Path of your project in Eclipse

- Here is what I got in Google search – click the first link
“MySQL::Download Connector/J”



Download MySQL JDBC driver and add it to Java Build Path of your project in Eclipse

- Download the zip file (the 2nd Download link)

The screenshot shows the MySQL Connector/J download page. The browser address bar displays `https://dev.mysql.com/downloads/connector/j/`. The page has a blue header with the MySQL logo and navigation tabs: MySQL.com, Downloads (selected), Documentation, and Developer Zone. Below the header, there are links for Enterprise, Community, Yum Repository, APT Repository, SUSE Repository, Windows, and Archives. The main content area is titled "Download Connector/J" and describes the driver as the official JDBC driver for MySQL. It includes links to online documentation and installation instructions. A sidebar on the left lists various MySQL products, with "MySQL Connectors" expanded to show Connector/ODBC, Connector/Net, Connector/J (selected), Connector/Python, Connector/C++, and Connector/C. The main content area shows a section for "Generally Available (GA) Releases" with a dropdown menu set to "Platform Independent". Below this, there are two download options: a "Compressed TAR Archive" (5.1.38, 3.8M) and a "ZIP Archive" (5.1.38, 4.1M), each with a "Download" button. A small box on the right states that MySQL open source software is provided under the GPL License and that OEMs, ISVs, and VARs can purchase commercial licenses.

MySQL Connector/J is the official JDBC driver for MySQL.

Online Documentation:

- MySQL Connector/J Installation Instructions, Documentation and Change History

Please report any bugs or inconsistencies you observe to our [Bugs Database](#).
Thank you for your support!

Generally Available (GA) Releases

Connector/J 5.1.38

Select Platform:
Platform Independent

Looking for previous GA versions?

Platform Independent (Architecture Independent), Compressed TAR Archive (mysql-connector-java-5.1.38.tar.gz)	5.1.38	3.8M	Download
Platform Independent (Architecture Independent), ZIP Archive	5.1.38	4.1M	Download

MD5: 8f8e768a91338328f2ac5cd6b6683c88 | [Signature](#)

MySQL open source software is provided under the [GPL License](#).
OEMs, ISVs and VARs can purchase commercial licenses.

Download MySQL JDBC driver and add it to Java Build Path of your project in Eclipse

- On the page that follows, click on “**No thanks, just start my download**”

MySQL.com Downloads Documentation Developer Zone

Yum Repository APT Repository SUSE Repository Windows Archives

Begin Your Download - mysql-connector-java-5.1.38.zip

Login Now or Sign Up for a free account.

An Oracle Web Account provides you with the following advantages:

- Fast access to MySQL software downloads
- Download technical White Papers and Presentations
- Post messages in the MySQL Discussion Forums
- Report and track bugs in the MySQL bug system
- Comment in the MySQL Documentation

Login »
using my Oracle Web account

Sign Up »
for an Oracle Web account

MySQL.com is using Oracle SSO for authentication. If you already have an Oracle Web account, click the Login link. Otherwise, you can signup for a free account by clicking the Sign Up link and following the instructions.

No thanks, just start my download.

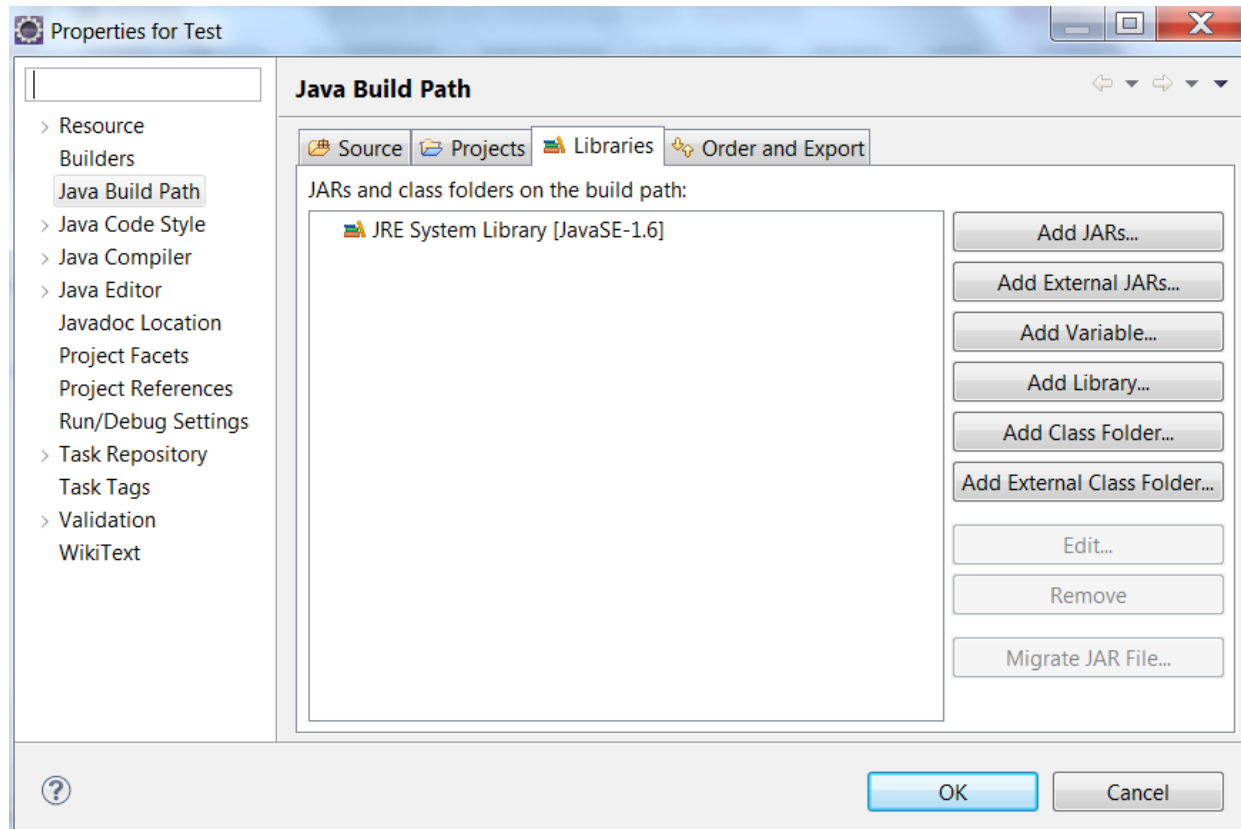
Download MySQL JDBC driver and add it to Java Build Path of your project in Eclipse

- Unzipped downloaded zip file contains one jar file, which is your jdbc driver
- You must add this jar file to Java Build Path of your project in Eclipse

Name	Date modified	Type	Size
docs	3/23/2016 5:06 PM	File folder	
src	3/23/2016 5:06 PM	File folder	
build.xml	12/2/2015 8:02 AM	XML Document	94 KB
CHANGES	12/2/2015 8:02 AM	File	233 KB
COPYING	12/2/2015 8:02 AM	File	18 KB
mysql-connector-java-5.1.38-bin.jar	12/2/2015 8:02 AM	Executable Jar File	961 KB
README	12/2/2015 8:02 AM	File	60 KB
README.txt	12/2/2015 8:02 AM	Text Document	63 KB

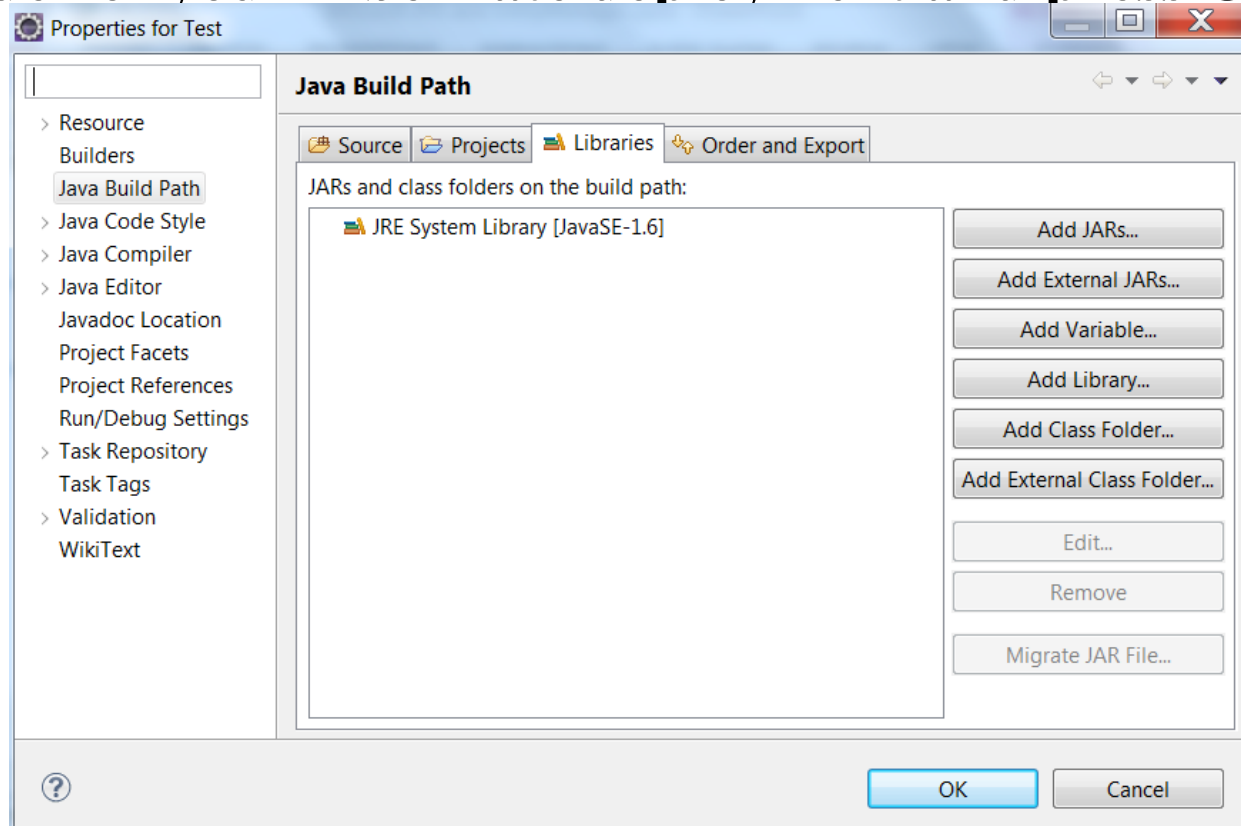
Adding Hibernate/JPA jars to your classpath in eclipse

- Select or Click on your **Java Project** or **Dynamic Web Project** name in Eclipse, and then select **Project->Properties** for your project
 - You will see something like this



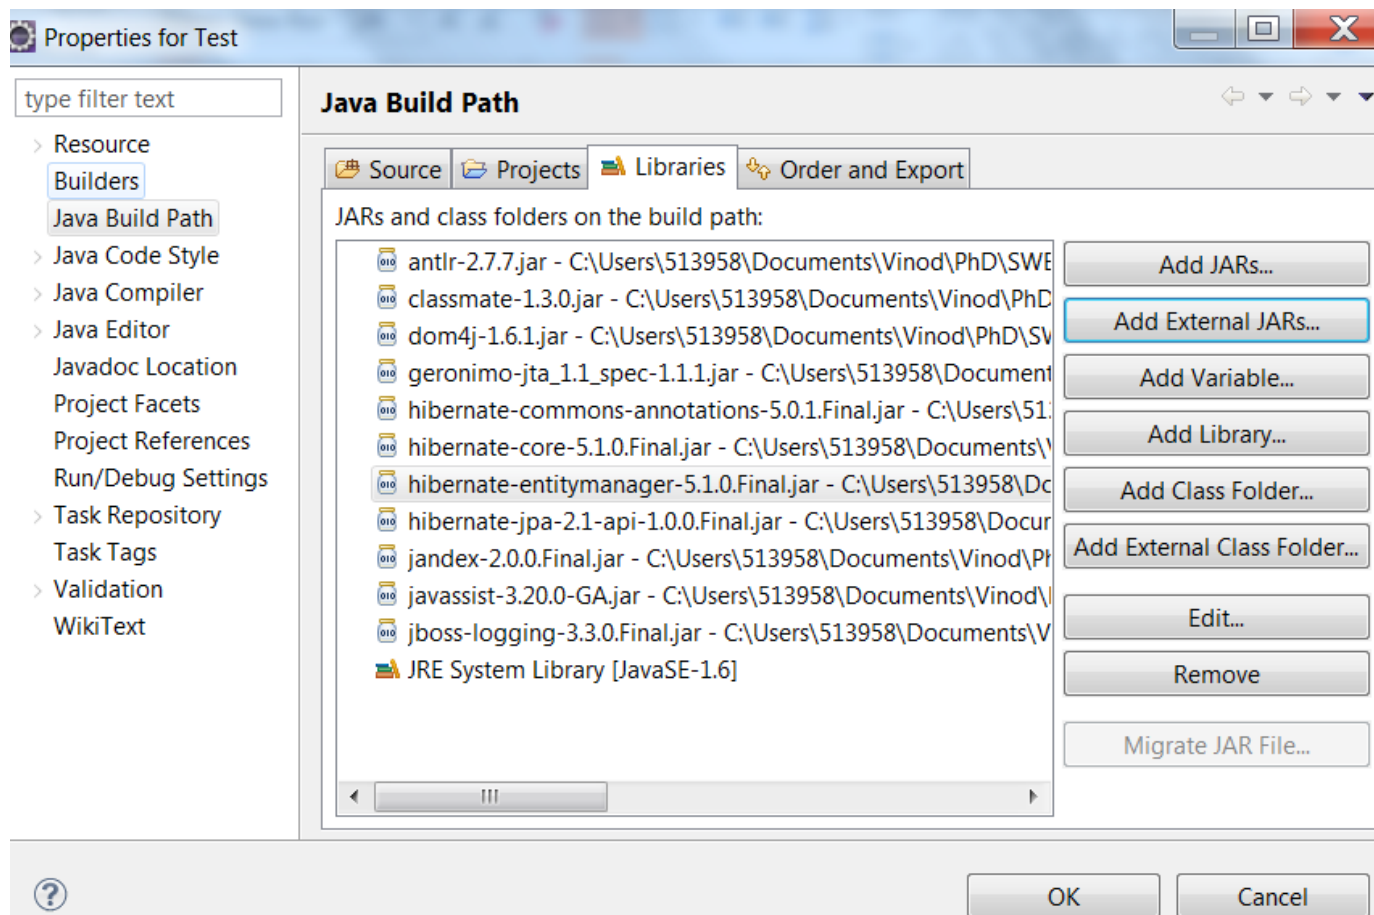
Adding Hibernate/JPA jars to your classpath in eclipse

- Now select **Java Build Path** (in the left panel) and the **Library** tab (as shown below)
- Then use **Add External Jars** link to add all jars in **lib/required** and **lib/jpa** folder of your hibernate deployment and press OK



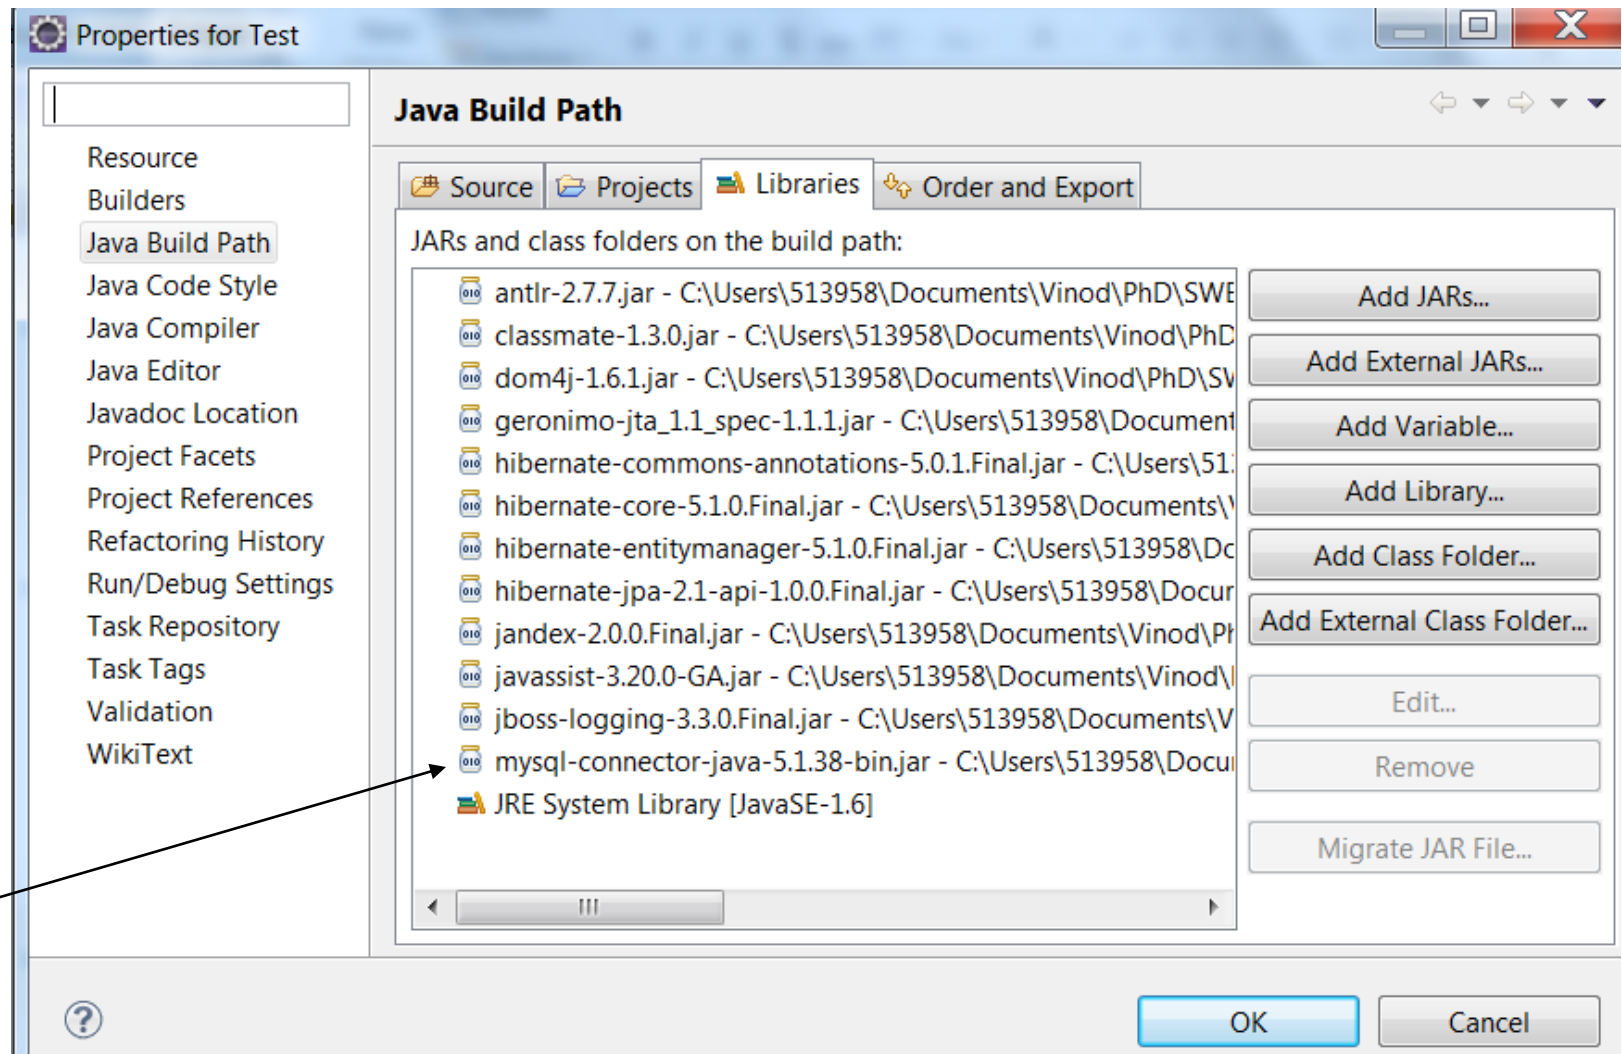
Adding Hibernate/JPA jars to your classpath in eclipse

- After adding the jar files, you should see something like as shown below – now your project's class path has **hibernate-core** jar and all its dependencies jars as well as jpa jars!



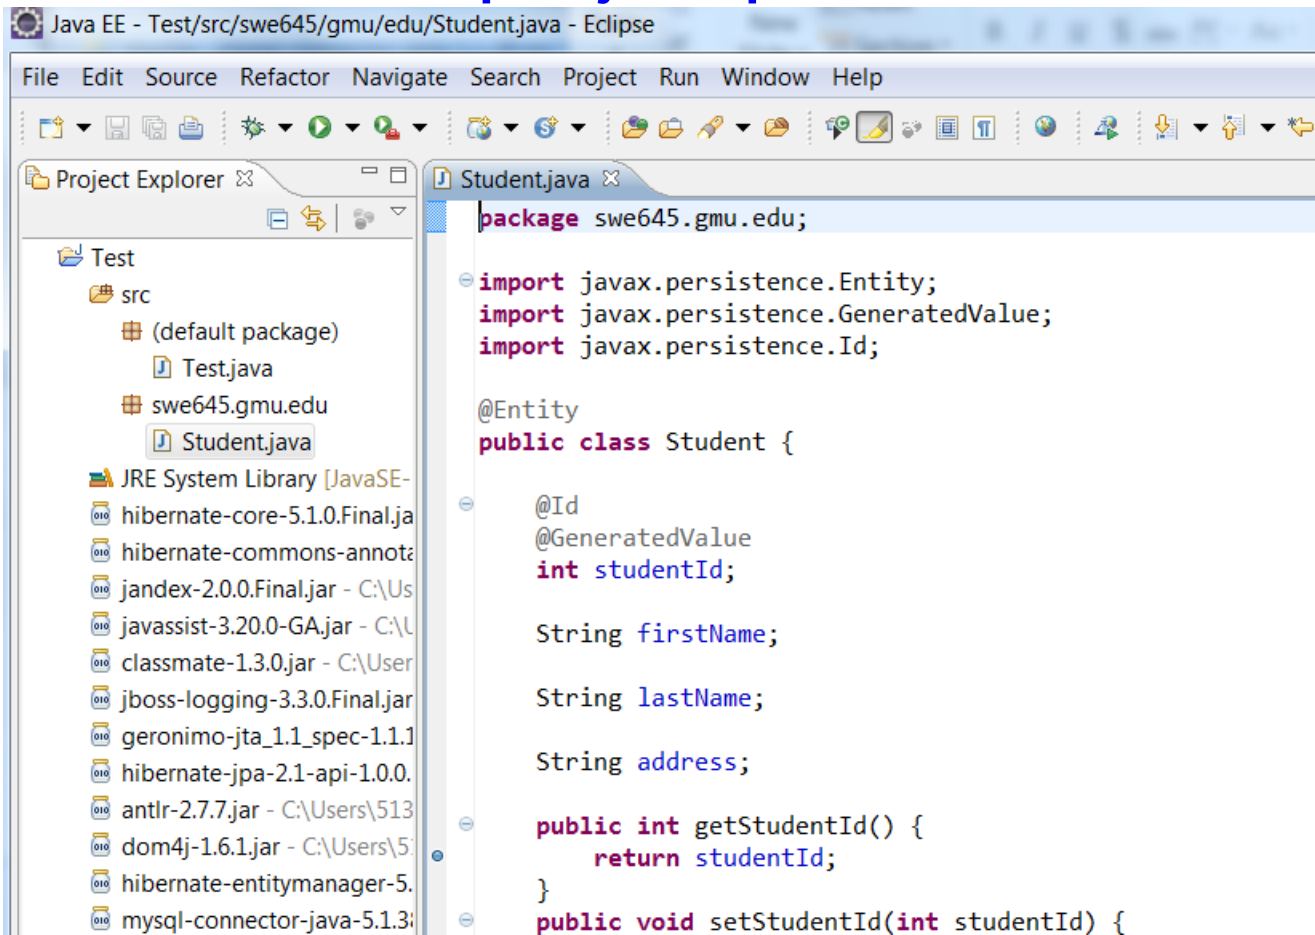
Adding Hibernate/JPA jars to your classpath in eclipse

- Repeat the same steps to add MySQL JDBC Driver jar



Your project is now ready to create JPA/Hibernate entity

- You can now create POJO and use JPA/hibernate annotations to create your JPA or hibernate entity
 - Make sure to use **import javax.persistence.***



The screenshot shows the Eclipse IDE interface. The title bar reads 'Java EE - Test/src/swe645/gmu/edu/Student.java - Eclipse'. The menu bar includes File, Edit, Source, Refactor, Navigate, Search, Project, Run, Window, and Help. The toolbar contains various icons for file operations and development tools. The Project Explorer on the left shows a project named 'Test' with a 'src' folder containing '(default package)', 'Test.java', and 'swe645.gmu.edu'. Under 'swe645.gmu.edu', the file 'Student.java' is selected. The main editor window displays the code for 'Student.java'. The code starts with the package declaration 'package swe645.gmu.edu;' followed by three imports: 'import javax.persistence.Entity;', 'import javax.persistence.GeneratedValue;', and 'import javax.persistence.Id;'. The class is annotated with '@Entity' and is a public class named 'Student'. It has three attributes: 'int studentId;' (annotated with '@Id' and '@GeneratedValue'), 'String firstName;', and 'String lastName;'. There is also a 'String address;' attribute. The class contains two methods: 'public int getStudentId() { return studentId; }' and 'public void setStudentId(int studentId) {'.

```
package swe645.gmu.edu;

import javax.persistence.Entity;
import javax.persistence.GeneratedValue;
import javax.persistence.Id;

@Entity
public class Student {

    @Id
    @GeneratedValue
    int studentId;

    String firstName;

    String lastName;


    String address;

    public int getStudentId() {
        return studentId;
    }

    public void setStudentId(int studentId) {
```

Create Hibernate JPA Configuration file: persistence.xml

- **R-Click** on your **Java Project** in Eclipse, select **New-File** and give it a name, **persistence.xml**, save it in **src/META** folder



Create Hibernate JPA Configuration file: persistence.xml

- Below is a template example of **persistence.xml**, which resides in the META-INF folder.

```
<persistence xmlns="http://java.sun.com/xml/ns/persistence"
            xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
            xsi:schemaLocation="http://java.sun.com/xml/ns/persistence http://java.sun.com/xml/ns/persistence/persistence_2_0.xsd"
            version="2.0">
  <persistence-unit name="sample">
    <provider>org.hibernate.ejb.HibernatePersistence</provider>
    <jta-data-source>java:/DefaultDS</jta-data-source>
    <mapping-file>ormap.xml</mapping-file>
    <jar-file>MyApp.jar</jar-file>
    <class>org.acme.Employee</class>
    <class>org.acme.Person</class>
    <class>org.acme.Address</class>
    <properties>
      <property name="hibernate.connection.driver_class">com.mysql.jdbc.Driver</property>
      <property name="hibernate.connection.password">XXXXXX</property>
      <property name="hibernate.connection.url">jdbc:mysql://<hostname>/<database></property>
      <property name="hibernate.connection.username">XXXXXX</property>
      <property name="hibernate.default_schema">XXXXXX</property>
      <property name="hibernate.dialect">org.hibernate.dialect.MySQLDialect</property>
    </properties>
  </persistence-unit>
</persistence>
```

Create Hibernate JPA Configuration file: persistence.xml

- Below is an example of **persistence.xml**, which resides in the **src/META-INF** folder. Here **swe645** is the MySQL database name created via Amazon RDS.

```
<?xml version="1.0" encoding="UTF-8"?>
- <persistence xsi:schemaLocation="http://xmlns.jcp.org/xml/ns/persistence
http://xmlns.jcp.org/xml/ns/persistence/persistence_2_1.xsd"
xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
xmlns="http://xmlns.jcp.org/xml/ns/persistence" version="2.1">
  - <persistence-unit transaction-type="RESOURCE_LOCAL" name="assign3">
    <provider>org.hibernate.ejb.HibernatePersistence</provider>
    <class>assignment.Student</class>
    - <properties>
      <property name="hibernate.dialect" value="org.hibernate.dialect.MySQLDialect"/>
      <property name="hibernate.hbm2ddl.auto" value="update"/>
      <property name="hibernate.show_sql" value="true"/>
      <property name="hibernate.connection.driver_class" value="com.mysql.jdbc.Driver"/>
      <property name="hibernate.connection.url" value="jdbc:mysql://127.0.0.1:3306/swe645"/>
      <property name="hibernate.connection.username" value="nsimon2"/>
      <property name="hibernate.connection.password" value="nopassword"/>
    </properties>
  </persistence-unit>
</persistence>
```


Create Hibernate JPA Configuration file: persistence.xml

- Below is an example of **persistence.xml** using GMU's Oracle database.

```
<?xml version="1.0" encoding="UTF-8"?>
- <persistence xsi:schemaLocation="http://xmlns.jcp.org/xml/ns/persistence
http://xmlns.jcp.org/xml/ns/persistence/persistence_2_1.xsd"
xsi:schemaLocation="http://www.w3.org/2001/XMLSchema-instance"
xmlns="http://xmlns.jcp.org/xml/ns/persistence" version="2.1">
  - <persistence-unit transaction-type="RESOURCE_LOCAL" name="SWE645HW4">
    <class>Student</class>
    - <properties>
      <property name="hibernate.connection.url"
        value="jdbc:oracle:thin:@apollo.ite.gmu.edu:1521:ite10g"/>
      <property name="hibernate.connection.driver_class" value="oracle.jdbc.driver.OracleDriver"/>
      <property name="hibernate.connection.username" value="sshres18"/>
      <property name="hibernate.connection.password" value="esoals"/>
      <property name="hibernate.archive.autodetection" value="class"/>
      <property name="hibernate.show_sql" value="true"/>
      <property name="hibernate.format_sql" value="true"/>
      <property name="hbm2ddl.auto" value="update"/>
    </properties>
  </persistence-unit>
</persistence>
```

JPA Providers

- **EclipseLink -- reference implementation**
 - Download the "EclipseLink Installer Zip" implementation from <http://www.eclipse.org/eclipselink/downloads/>
 - The download contains several jars. You may need the following jars:
 - eclipselink.jar
 - javax.persistence_*.jar

Summary

- **Java Persistence API**
 - **JPA Entities** -- How to Use Annotations to configure POJOs into JPA entities
- **Entity Manager**
 - Entity Manager Methods to interact with persistence storage
 - Persistence Unit/Persistence Context
- **Examples**

JPA Entity Relationships

Topics in this section

- **Intro to JPA Entity Relationships**
 - Cardinality
 - Unidirectional vs. bidirectional
- **Seven Entity Relationships**
- **Ordered List-based Relationships**
- **Cascading**
- **Lazy vs. Eager Loading**

Entity Relationships

- **Entity Relationships** are used **to model real-world business concepts**
- **Entity beans** must be capable of **forming relationships**
- **For instance:**
 - An employee has an address or
 - An employee has multiple phone numbers

Entity Relationships

- **Seven types of relationships** can exist between entity beans
- **There are 4 types of cardinality:**
 - one-to-one,
 - one-to-many,
 - many-to-one, and
 - many-to-many
- **Each relationship** can be either **unidirectional** or **bidirectional**

The Seven Relationship Types

- **One-to-one unidirectional**
- **One-to-one bidirectional**
- **One-to-many unidirectional**
- **One-to-many bidirectional**
- **Many-to-one unidirectional**
- **Many-to-many unidirectional**
- **Many-to-many bidirectional**

JPA Support for Entity Relationship

- The relationships between different entities are defined using **@OneToOne**, **@OneToMany**, **@ManyToOne**, and **@ManyToMany** annotations
 - A **unidirectional** relationship requires the owning side to specify the annotation.
 - A **bidirectional** relationship also requires the non-owning side to refer to its **owning side** by use of the **mappedBy** element of the **OneToOne**, **OneToMany**, or **ManyToMany** annotation.



One-to-One Unidirectional Relationship

- **One-to-One relationship** refers to a relationship where **one entity is related to exactly one other entity**
 - For example, relationship between **Customer** and **Address** entities
 - **Each customer has exactly one address**, and each address belongs to one customer
- **Unidirectional** means **only one entity** has **access method** to get the other object

One-to-One Unidirectional Relationship

- In the next example, **Address_ID** is a **join column** in both tables
- **getAddress()** in **Customer** table provides access to **Address** entity

One-to-One Unidirectional Relationship



Bean Class

- A **unidirectional relationship** requires the **owning side to specify the annotation**

@Entity

```
public class Customer implements Serializable{  
    private Address address;
```

@OneToOne



@JoinColumn(name="Address_ID")

```
public Address getAddress() {  
    return address;  
}  
public void setAddress(Address address) {  
    this.address = address;  
}
```

One-to-One Bi-directional Relationship

- **Both entities have methods to access the other object**
- **The following example, creates a join column CC_Number in both tables**

One-to-One Bi-directional Relationship



One-to-One Bi-directional Relationship

- In bi-directional relation, need to **specify** that the **mapping** is **done** by **one entity**
- A **bidirectional** relationship requires the **non-owning side** to **refer** to its **owning side** by use of the **mappedBy** element of the **OneToOne**, **OneToMany**, or **ManyToMany** annotation.

Bean Classes

@Entity

```
public class Customer implements Serializable{  
    private CreditCard creditCard;
```

@OneToOne

@JoinColumn (name="CC_Number")

```
    public CreditCard getCreditCard() {return creditCard;}  
    public void setCreditCard(CreditCard cc)  
    { creditCard = cc; }
```

@Entity

```
public class CreditCard implements Serializable {  
    private Customer cust;
```

@OneToOne (mappedBy="**creditCard**")



```
    Public Customer getCustomer() { returncust; }
```

- In bi-directional relation, need to specify, in the non owning side, that the mapping is done by the owning side entity,.

One-to-Many Unidirectional Relationship

- **One-to-Many relationship** refers to a relationship where **one entity** is **related to more than one other entity**
 - For example, relationship between **Customer** and **Phone** entities
 - **A customer has many phones**
- Relationship is **unidirectional** means **only customer entity** has **access method** to get the phone objects

One-to-Many Unidirectional Relationship



- A Customer has many Phones

Bean Class

`@Entity`

```
public class Customer implements Serializable{  
    private Collection<Phone> phoneNumbers = ... ..
```

`@OneToMany`

`@JoinColumn(name="Person_ID")`



```
public Collection<Phone> getPhones() {  
    return phones;  
}
```

```
public void setPhone(Collection<Phone> phones) {  
    this.phones = phones;  
}
```

One-to-Many Relationship With Join Table

- The Hibernate, by default, creates a **separate mapping table** (or join table) for **one to many relationships**
 - Default join table name is concatenation of two entity names
 - Default join column names are `tablename+columnname` for both tables
- You can override default names for join columns and join tables created by hibernate

One-to-Many Relationship With Join Table



- The Hibernate, by default, creates a separate mapping table (or join table) for one to many relationships.

Bean Class

`@Entity`

```
public class Customer implements Serializable{  
    private Collection<Phone> phoneNumbers = ... ..
```

`@OneToMany`


```
    @JoinTable (name="Cust_Phone_Rel",  
        joinColumns={@JoinColumn (name="Person_ID") },  
        inverseJoinColumns={@JoinColumn (name="Phone_ID") })  
    public Collection<Phone> getPhones () {  
        return phones;  
    }  
  
    public void setPhone(Collection<Phone> phones) {  
        this.phones = phones;  
    }
```

- You can override default names for join columns and join tables created by hibernate
- Default join table name is concatenation of two entity names
- Default join column names are tablename+columnname for both tables

One-to-Many Bi-directional Relationship without a separate Join table

- Another way to represent one-to-many relationship is to have **the many side** of the relationship (e.g., **Reservation**) a reference to one-side of the relationship (e.g., **Cruise**).
 - That is **Reservation** has **CruiseID** as column instead of creating a separate mapping table.
- To accomplish this, **use mappedBy** of **OneToMany** relationship
 - to say that **Cruise** instance of **Reservation** class has to do the **mapping**, as shown on the next page.
 - This creates only two tables **Cruise** and **Reservation**.

One-to-Many Bi-directional Relationship




- Another way to represent one-to-many relationship is to have the many side of the relationship (e.g., Reservation) a reference to one-side of the relationship (e.g., Course). That is Reservation has CruiseID as column instead of creating a separate mapping table.
- To accomplish this, use `mappedBy` of `OneToMany` relationship to say that Cruise instance of Reservation class has to do the mapping, as shown on the next page. This creates only two tables Cruise and Reservation.

Bean Classes

```
@Entity
public class Reservation implements Serializable{
    private Cruise cruise;

    @ManyToOne
    @JoinColumn(name="Cruise_Id")
    public Cruise getCruise()
    public void setCruise(...)
```



```
@Entity
public class Cruise implements Serializable{
    private Collection<Reservation> reservations;

    @OneToMany(mappedBy="cruise")
    public Collection<Reservation> getReservations()
    { .... }
```

mappedBy used to tell that mapping is done by Cruise instance in Reservation and this creates a Cruise_Id column in Reservation table instead of creating a separate join table.

Ordered List-Based Relationships

```
@Entity
public class CruiseShip implements Serializable{
    private Collection<Cabin> cabins;

    @OneToMany
    @OrderBy("roomClass asc, roomNumber desc")
    @JoinColumn(name="Ship_Id")
    public Collection<Cabin> getCabins() {
        return this.cabins;
    }
}
```


Cascading

- When you **perform** an **EntityManager operation** on an **entity bean** instance, you can **automatically** have the **same operation** performed **on any relationship** properties the entity may have. This is called **cascading**.
- **Cascading** can be **applied** to a variety of **EntityManager operations**, including
 - **`persist()`, `merge()`, `remove()`, and `refresh()`.**
- This feature is enabled by setting the **`javax.persistence.CascadeType`** of the relationship annotation's ***cascade* attribute**.

Cascading

- The **CascadeType** is defined as a Java enumeration

```
public enum CascadeType
{
    ALL, PERSIST,
    MERGE, REMOVE,
    REFRESH
}
```

- The **ALL** value represents **all** of the **cascade** operations.
- The remaining values represent individual cascade operations.

Cascading

- For example, persisting a new **Employee** entity with a new **address** and **phone number**
- All you have to do is **wire the object**, and the entity manager can **automatically create the employee and its related entities**, all in one `persist()` method call

```
Employee employee = new Employee();  
employee.setAddress(new Address());  
employee.getPhoneNumbers().add(new Phone());  
  
// create them all in one entity manager invocation  
entityManager.persist(employee);
```

Cascading

Used in relationship annotation

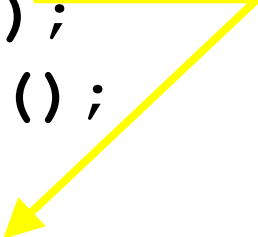
`@OneToMany(cascade={CascadeType.ALL})`

```
public enum CascadeType {  
    ALL, PERSIST, MERGE, REMOVE, REFRESH  
}
```

Persist

```
Customer cust = new Customer();  
Address address = new Address();  
cust.setAddress(address);  
entityManager.persist(cust);
```

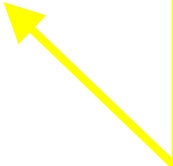
Customer and
Address are
persisted.



Cascading

Merge

```
Cust.setName("William");  
Cust.getAddress().setCity("Boston");  
entityManager.merge(cust);  
Phone phone = new Phone();  
Phone.setNumber("123456789");  
cust.getPhoneNumbers().add(phone);  
entityManager.merge(cust);
```



**Customer and
Address are
merged**

Cascading

Remove

```
Customer cust =  
entityManager.find(Cstomer.class, 1);  
entityManager.remove(cust);
```

Refresh

```
Customer cust =  
entityManager.find(Cstomer.class, 1);  
entityManager.refresh(cust);
```

Lazy vs. Eager Loading

- The **FetchType.EAGER** annotation may be specified on an entity **to eagerly load the data** from the database.
- The **FetchType.LAZY** annotation may be specified as **a hint** that **the data** should be **fetches lazily** when it is first accessed.

Lazy vs. Eager Loading

- Consider a **relationship** between **University** and **Student** entities
 - The University entity might have some basic properties such as id, name, address, etc. as well as a property called students:

```
public class University {  
    private String id;  
    private String name;  
    private String address;  
    private List<Student> students;  
  
    // setters and getters  
}
```


Lazy vs. Eager Loading

- **When a University entity is loaded from the database, JPA loads its id, name, and address fields for you.**
- **There are two options for students:**
 - to load it together with the rest of the fields (i.e. **eagerly**) or
 - to load it on-demand (i.e. **lazily**) when you call the university's `getStudents()` method.

Lazy vs. Eager Loading

- When a **university** has **many students** it is **not efficient** to **load all** of its **students** with it when they are not needed.
- So in such cases, you can declare that you want **students to be loaded when** they are **actually needed**. This is called **lazy loading**.

Lazy vs. Eager Loading

Relationship Type	Default Fetch Behavior	Number of Entity Retrieved
One-To-One	Eager	Single Entity
One-To-Many	Lazy	Collection of Entities
Many-To-One	Eager	Single Entity
Many-To-Many	Lazy	Collection of Entities

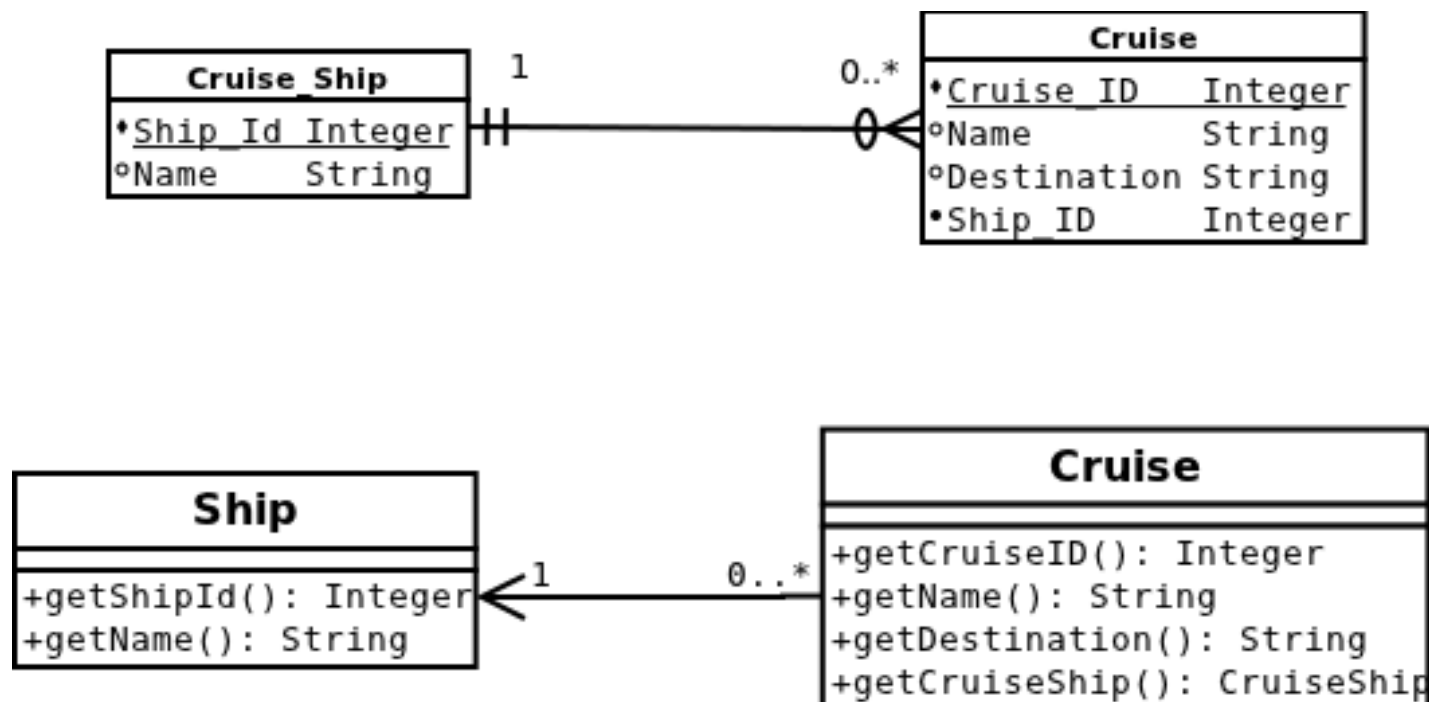
Lazy vs. Eager Loading

`@Entity`

```
public class Customer implements Serializable{  
  
    private Collection<Phone> phones;  
  
    @OneToMany (cascade={ CascadeType.ALL} ,  
                fetch=FetchType.EAGER)  
    @JoinColumn (name="customer_id")  
    public Collection<Phone> getPhones() {  
        return phones;  
    }  
}
```

Optional Topics

Many-to-One Unidirectional



Bean Class

@Entity

```
public class Cruise implements Serializable{  
    private CruiseShip ship;
```

@ManyToOne

@JoinColumn(name="Ship_Id")



```
public CruiseShip getCruiseShip() {  
    return ship;  
}
```

```
public void setCruiseShip(CruiseShip ship) {  
    this.ship = ship;  
}
```

Many-to-Many Bidirectional Relationship

- **Bidirectional** means both entities have methods to access the other objects(s)
- **Many-to-Many relationship** has to have a separate mapping (i.e., join) table
- **For example:**
 - Res_Cust_Rel table
 - where names of table and columns can be changed using **joinColumns** and **inverseJoinColumns** property as shown on the next page.

Many-to-Many Bidirectional Relationship

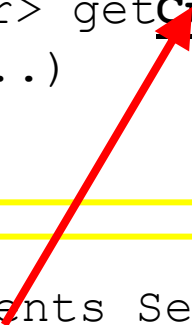


Many-to-Many Bidirectional Relationship

- For **@ManyToMany** annotations in both tables, hibernate, by default, will create two mapping tables
 - Mapping Reservation to Customer and
 - Mapping Customer to Reservation
- This is **redundant** and **can be overridden**.
- That is, create only **one mapping table (i.e., join table)** by using **mappedBy** property in **one method** as shown on the next page.

Bean Classes

```
@Entity
public class Reservation implements Serializable{
    @ManyToMany
    @JoinTable (name="Res_Cust_Rel",
        joinColumns{@JoinColumn (name="Reservation_ID")},
        inverseJoinColumns={@JoinColumn (name="Customer_ID") })
    public Collection<Customer> getCustomers()
    public void setCustomer(...)
```





```
@Entity
public class Customer implements Serializable{
    @ManyToMany (mappedBy="customers")
    public Collection<Reservation> get Reservations()
    public void setReservations(...)
```

Here we say that mapping is done by the other entity which creates a join table, so don't need to create another mapping.

Many-to-Many Unidirectional Relationship

- **Only one entity has access method to get the other objects**
 - Reservation can get all its related Cabins, but
 - Cabin does not know who made Reservation

Many-to-Many Unidirectional Relationship



No get method in Cabin to access Reservation

Bean Class

```
@Entity
public class Reservation implements Serializable{
    private Collection<Cabin> cabins;

    @ManyToMany
    @JoinTable(name="Res_Cabin_Rel",
        joinColumns={@JoinColumn(name="Reservation_Id")},
        inverseJoinColumns={@JoinColumn(name="Cabin_Id")})
    public Collection<Cabin> getCabins()

    public void setCabins(...)
```

JDBC Data Source

A **basic DataSource** implementation

Standard connection object that is **not pooled** or **used** in a distributed transaction.

A **DataSource** class that **supports connection pooling**

Connection objects that **participate** in connection pooling and will be recycled.

A **DataSource** class that **supports distributed transaction**

Connections objects that can be used in a distributed transaction.

Obtain Persistence Context inside EJB

A reference to the Persistence Context can be injected into the EntityManager.

```
@Stateless
public class StatelessEJB implements
StatelessRemote {

    @PersistenceContext(unitName="swe645Unit")
    private EntityManager entityManager;
```


Tips for Designing Entity Beans

Read an application description and underline the **Nouns**. These nouns can be **representative** of the **Entities (tables)**.

Determine the **relationship between the Entities**.

How do they relate?

Are they one to many? Many to many? (Discussed later.)

System Description (example)

A new university is opening. There are 3 locations, VA, MD, and DC. The people who attend this university are Students and Employees. Employees and Students have contact information (name, phone number, address, email, and date of birth). Employees have a title, department, hire date, employment status, and a termination date. Students have a Major, start date, GPA, graduation date, and status. Employees can be Students. Students can register for multiple courses. Each course can have many students. Courses are assigned to a specific room, in a building, for a location. Each location has multiple buildings. Each building has multiple rooms.